

Machine Learning Lab II

- PS1
- NUMPY

PS1 Q&A

Numpy Tutorial

Overview

- Numpy: basic objects, methods, functions
- Numpy: linear algebra
- Numpy: random
- Matplotlib: 2D plots
- Matplotlib: 3D plots
- Scipy vs Numpy
- Discuss assignment 1

Numpy

- Fundamental package for working with N-dimensional array objects (vector, matrix, tensor, ...)
- Numpy arrays are a fundamental data type for some other packages to use
- Numpy has many specialized modules and functions:
 - `numpy.linalg` (Linear algebra)
 - `numpy.random` (Random sampling)
 - `numpy.fft` (Discrete Fourier transform)
 - sorting/searching/counting
 - math functions
 - `numpy.testing` (unit test support)

Declaring a Numpy array

- Each Numpy array has some attributes:
 - shape (a tuple of the size in each dimension)
 - dtype (data type of entries)
 - size (total # of entries)
 - ndim (# of dimensions)
 - T (transpose)

What can you do?

- Add two arrays
- Add all entries in one array
- Multiply two arrays (1D, 2D)
- Take the exponential of each element in an array
- Multiply an array by a scalar
- Get the minimum element of an array
- Print a few elements of an array
- Print a single column or row of an array
- Multiply two arrays via matrix multiplication

Array broadcasting

0	0	0
10	10	10

 +

0	1	2
---	---	---

 =

0	1	2
10	11	12

0	0	0
10	10	10

 +

0	1	2
0	1	2

 =

0	1	2
10	11	12

0
10

 +

0	1	2
---	---	---

 =

0	1	2
10	11	12

Iterating over an array

- Iteration over all elements of array:
 - for element in A.flat
- Iteration over multidimensional arrays is done on slices in the first dimension:
 - for row in A
- Alternatively, could access entries through indices:
 - for i in range(A.shape[0]):
 - for j in range(A.shape[1]):

Reshaping an array

- Use reshape to modify the dimensions of an array while leaving the total number of elements the same
 - `A = np.arange(8)`
 - `A.reshape(2,4)`
 - gives `[[0,1,2,3],[4,5,6,7]]`
- Use resize to remove elements or append 0's in place
 - (size can change under some circumstances*)
 - `resize(2,3)`
- Use resize to return a copy with removed elements or repeated copies
 - `b = resize(a,(2,4))`

Numpy: Linear Algebra

name	explanation
<code>dot(a,b)</code>	dot product of two arrays
<code>kron(a,b)</code>	Kronecker product
<code>linalg.norm(x)</code>	matrix or vector norm
<code>linalg.cond(x)</code>	condition number
<code>linalg.solve(A,b)</code>	solve linear system $Ax=b$
<code>linalg.inv(A)</code>	inverse of A
<code>linalg.pinv(A)</code>	pseudo-inverse of A
<code>linalg.eig(A)</code>	eigenvalues/vectors of square A
<code>linalg.eigvals(A)</code>	eigenvalues of general A
<code>trace(A)</code>	trace (diagonal sum)
<code>linalg.svd(A)</code>	singular value decomposition

<http://docs.scipy.org/doc/numpy/reference/routines.linalg.html>

Numpy: Random

- `x = np.random.randn(50)`
- `y = 3.5*x+2+np.random.randn(50)*0.3`
- If you run this, you'll get different numbers each time, so you might want to use `np.random.seed(*)` to reproduce a random experiment

Numpy: Random

name	explanation
<code>rand(n0,n1,...)</code>	ndarray of random values from uniform [0,1]
<code>randn(n0,n1,...)</code>	random standard normal
<code>randint(lo, [hi, size])</code>	random integers [lo, hi)
<code>shuffle(seq)</code>	shuffle sequence randomly
<code>choice(seq,[size,replace,p])</code>	sample k items from a 1D array with or without replacement
<code>chisquare(df,[size])</code>	sample from Chi-squared distribution with df degrees of freedom
<code>exponential([scale,size])</code>	sample from exponential distribution

<http://docs.scipy.org/doc/numpy/reference/routines.random.html>

- Matplotlib is the 2D Python plotting library
- We'll mostly use `matplotlib.pyplot`
- There are tons of options, so consult the documentation:
 - <http://matplotlib.org/users/beginner.html>
- `matplotlib.pyplot` can do many types of visualizations including:
 - Histograms, bar charts (using `hist`)
 - Error bars on plots, box plots (using `boxplot`, `errorbar`)
 - Scatterplots (using `scatter`)
 - Line plots (using `plot`)
 - Contour maps (using `contour` or `tricontour`)
 - Images (matrix to image) (using `imshow`)
 - Stream plots which show derivatives at many locations (`streamplot`)
 - Pie charts, polar charts (using `pie`, `polar`)

Matplotlib: 2D plots

- How do we show two curves on the same plot?
 - `import numpy as np`
 - `import matplotlib.pyplot as plt`
 - `x = np.linspace(0,np.pi,100)`
 - `y = np.sin(x)`
 - `plt.plot(x,y)`
 - `plt.show()`
- More examples: <http://matplotlib.org/gallery.html>
- Documentation: http://matplotlib.org/api/pyplot_api.html

Scipy vs. Numpy

- Scipy is a library that can work with Numpy arrays, but can achieve better performance and has some more specialized libraries
 - linear algebra (scipy.linalg uses BLAS/LAPACK)
 - statistics (scipy.stats has hypothesis tests, correlation analysis)
 - optimization (scipy.optimize has multiple solvers, gradient checks, simulated annealing)
 - sparse matrices (scipy.sparse supports sparse linear algebra, graph analysis, multiple sparse matrix formats)
 - signal processing (scipy.signal has convolutions, wavelets, splines, filters)

<http://docs.scipy.org/doc/scipy/reference/>

References

- Stanford CME 193 By Eileen Martin
- <https://docs.scipy.org/doc/numpy/index.html>